

Fraud Detection using Logistic Regression

December 8, 2022

```
[1]: import pandas as pd
```

```
df = pd.read_csv('data/card_data.csv')
```

```
[2]: df.head()
```

```
[2]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	

	V8	V9	...	V21	V22	V23	V24	V25	\
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]

```
[3]: df['Class'].value_counts()
```

```
[3]:
```

0	284315
1	492

Name: Class, dtype: int64

```
[4]: df = df.sample(frac = 1, random_state = 1)
df = df.reset_index(drop = True)
```

```
df.head()
```

```
[4]:      Time      V1      V2      V3      V4      V5      V6 \
0  119907.0 -0.611712 -0.769705 -0.149759 -0.224877  2.028577 -2.019887
1   78340.0 -0.814682  1.319219  1.329415  0.027273 -0.284871 -0.653985
2   82382.0 -0.318193  1.118618  0.969864 -0.127052  0.569563 -0.532484
3   31717.0 -1.328271  1.018378  1.775426 -1.574193 -0.117696 -0.457733
4   80923.0  1.276712  0.617120 -0.578014  0.879173  0.061706 -1.472002

      V7      V8      V9  ...      V21      V22      V23      V24 \
0  0.292491 -0.523020  0.358468  ... -0.075208  0.045536  0.380739  0.023440
1  0.321552  0.435975 -0.704298  ... -0.128619 -0.368565  0.090660  0.401147
2  0.706252 -0.064966 -0.463271  ... -0.305402 -0.774704 -0.123884 -0.495687
3  0.681867 -0.031641  0.383872  ... -0.220815 -0.419013 -0.239197  0.009967
4  0.373692 -0.287204 -0.084482  ... -0.160161 -0.430404 -0.076738  0.258708

      V25      V26      V27      V28  Amount  Class
0 -2.220686 -0.201146  0.066501  0.221180     1.79      0
1 -0.261034  0.080621  0.162427  0.059456     1.98      0
2 -0.018148  0.121679  0.249050  0.092516     0.89      0
3  0.232829  0.814177  0.098797 -0.004273    15.98      0
4  0.552170  0.370701 -0.034255  0.041709     0.76      0
```

```
[5 rows x 31 columns]
```

```
[5]: as_np = df.to_numpy()
      index = int(len(as_np) * .92)

      X_train, y_train = as_np[:index, :-1], as_np[:index, -1]
      X_test, y_test = as_np[index:, :-1], as_np[index:, -1]

      (X_train.shape, y_train.shape), (X_test.shape, y_test.shape)
```

```
[5]: (((262022, 30), (262022,)), ((22785, 30), (22785,)))
```

```
[6]: from sklearn.preprocessing import StandardScaler

      scaler = StandardScaler().fit(X_train)

      X_train = scaler.transform(X_train)
      X_test = scaler.transform(X_test)

      X_test[0]
```

```
[6]: array([ 0.14097956,  0.53955733, -1.15153973, -0.47041404,  0.57191953,
          -0.85362208, -0.27419086, -0.03159233, -0.25697594,  2.43387034,
          -0.75622807, -0.03956163, -1.77401948,  2.41251471,  1.26340856,
```

```
-0.24099657, 0.16722599, 0.24463032, 0.56588687, -0.53397987,  
1.02654979, 0.4423462, 0.47928779, -0.49696321, -0.14435544,  
-0.64466426, 0.08065479, -0.24695714, 0.11458447, 1.64640304]]
```

```
[7]: from sklearn.linear_model import LogisticRegression  
  
model = LogisticRegression().fit(X_train, y_train)  
test_predictions = model.predict(X_test)  
  
pd.value_counts(test_predictions)
```

```
[7]: 0.0    22757  
     1.0     28  
     dtype: int64
```

```
[8]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay  
  
cm = confusion_matrix(y_test, test_predictions, labels = [0, 1])  
disp = ConfusionMatrixDisplay(confusion_matrix = cm,  
                              display_labels = ['Not Fraud', 'Fraud'])  
disp.plot()
```

```
[8]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x173812d02b0>
```



```
[9]: tn, fp, fn, tp = cm.ravel()

s = '''
True negatives: {0}
False Positives: {1}
False Negatives: {2}
True Positives: {3}'''.format(tn, fp, fn, tp)

print(s)
```

```
True negatives: 22741
False Positives: 8
False Negatives: 16
True Positives: 20
```

```
[10]: # accuracy = (tp+tn)/(tp+tn+fp+fn)

def accuracy(tn, fp, fn, tp):
    return ((tp+tn)/(tp+tn+fp+fn))

"Accuracy: {0}".format(accuracy(tn, fp, fn, tp))
```

```
[10]: 'Accuracy: 0.9989466754443713'
```

```
[11]: # recall = sensitivity = true positive rate = tp/(tp+fn)

def tpr(tn, fp, fn, tp):
    return (tp / (tp + fn))

"True positive Rate: {0}".format(tpr(tn, fp, fn, tp))
```

```
[11]: 'True positive Rate: 0.5555555555555556'
```

```
[12]: # false negative rate = fn/(tp+fn)

def fpr(tn, fp, fn, tp):
    return (fn / (tp + fn))

"False Negative Rate: {0}".format(fpr(tn, fp, fn, tp))
```

```
[12]: 'False Negative Rate: 0.4444444444444444'
```

```
[13]: # specificity = true negative rate = tn/(tn+fp)

def tnr(tn, fp, fn, tp):
    return (tn / (tn + fp))
```

```
"True Negative Rate: {0}".format(tnr(tn, fp, fn, tp))
```

```
[13]: 'True Negative Rate: 0.9996483361906018'
```

```
[14]: # false positive rate = fp/(tn+fp)
```

```
def fpr(tn, fp, fn, tp):  
    return (fp / (tn + fp))
```

```
"False Positive Rate: {0}".format(fpr(tn, fp, fn, tp))
```

```
[14]: 'False Positive Rate: 0.00035166380939821533'
```

```
[15]: # precision = positive predictive value = tp/(tp+fp)
```

```
def ppv(tn, fp, fn, tp):  
    return (tp / (tp + fp))
```

```
"Positive Predictive Value: {0}".format(ppv(tn, fp, fn, tp))
```

```
[15]: 'Positive Predictive Value: 0.7142857142857143'
```

```
[16]: # negative predictive value = tn/(tn+fn)
```

```
def npv(tn, fp, fn, tp):  
    return (tn / (tn + fn))
```

```
"Negative Predictive Value: {0}".format(npv(tn, fp, fn, tp))
```

```
[16]: 'Negative Predictive Value: 0.9992969196291251'
```

```
[17]: # balanced accuracy = (tpr+tnr)/2
```

```
def balanced_accuracy(tn, fp, fn, tp):  
    return (tpr(tn, fp, fn, tp) + tnr(tn, fp, fn, tp)) / 2
```

```
"Balanced Accuracy: {0}".format(balanced_accuracy(tn, fp, fn, tp))
```

```
[17]: 'Balanced Accuracy: 0.7776019458730787'
```

```
[18]: # f1 score = 2 x (precision x recall)/(precision + recall)
```

```
def f1(tn, fp, fn, tp):  
    p = ppv(tn, fp, fn, tp)  
    r = tpr(tn, fp, fn, tp)  
    return (2*p*r)/(p+r)
```

```
"F1 Score: {0}".format(f1(tn, fp, fn, tp))
```

```
[18]: 'F1 Score: 0.6250000000000001'
```

```
[19]: probabilities = model.predict_proba(X_test)[: , 1]

probabilities
```

```
[19]: array([0.00012139, 0.0003532 , 0.00030147, ..., 0.00473659, 0.00017273,
          0.00171865])
```

```
[20]: pd.value_counts(probabilities > 0.5)
```

```
[20]: False    22757
      True      28
      dtype: int64
```

```
[21]: import numpy as np

thresholds = np.linspace(0, 1, num = 2000).astype(np.float16)

thresholds
```

```
[21]: array([0.000e+00, 5.002e-04, 1.000e-03, ..., 9.990e-01, 9.995e-01,
          1.000e+00], dtype=float16)
```

```
[22]: all_predictions = np.array([(probabilities > t).astype(int) for t in
      ↪ thresholds])

all_predictions.shape
```

```
[22]: (2000, 22785)
```

```
[23]: all_predictions[-4]
```

```
[23]: array([0, 0, 0, ..., 0, 0, 0])
```

```
[24]: pd.value_counts(all_predictions[-4])
```

```
[24]: 0    22776
      1      9
      dtype: int64
```

```
[25]: confusion_matrices = [confusion_matrix(y_test, predictions) for predictions in
      ↪ all_predictions]
      tn_fp_fn_tps = [cm.ravel() for cm in confusion_matrices]

      tprs = [tpr(tn, fp, fn, tp) for tn, fp, fn, tp in tn_fp_fn_tps]
      fprs = [fpr(tn, fp, fn, tp) for tn, fp, fn, tp in tn_fp_fn_tps]
```

```
[26]: import plotly.express as px

px.scatter(x = fprs, y = tprs, color = thresholds, labels = dict(x = 'False_
↳Positive Rate', y = 'True Positive Rate', color = 'Threshold'), title = 'ROC_
↳Curve')
```

```
[27]: from sklearn.metrics import auc

auc(fprs, tprs)
```

```
[27]: 0.9810064911278153
```

```
[28]: # classification report

from sklearn.metrics import classification_report

print(classification_report(y_test, test_predictions, labels = [0, 1],
↳target_names = ['Not Fraud', 'Fraud']))
```

	precision	recall	f1-score	support
Not Fraud	1.00	1.00	1.00	22749
Fraud	0.71	0.56	0.63	36
accuracy			1.00	22785
macro avg	0.86	0.78	0.81	22785
weighted avg	1.00	1.00	1.00	22785

```
[29]: # matthews correlation coefficient = (tp x tn) - (fp x fn) / sqrt((tp + fp) x
↳(tp + fn) x (tn + fp) x (tn + fn))

import math as mt

def mcc(tn, fp, fn, tp):
    return (((tp*tn) - (fp*fn))/mt.sqrt((tp + fp)*(tp + fn)*(tn + fp)*(tn +
↳fn)))

"Matthews Correlation Coefficient: {0}".format(mcc(tn, fp, fn, tp))
```

```
[29]: 'Matthews Correlation Coefficient: 0.6294313746803477'
```

```
[ ]:
```